

# Hadoop version 3.x

# Agenda

- Overview of new features in Hadoop version 3
- Understanding the basics of erasure coding
- Understanding the features of intra data node balancer

# Salient features Hadoop version 3

- **Stable version 3.0.0 available as on 13<sup>th</sup> December 2017**
- **Minimum required Java version is Java 8**
- **Support for erasure coding in HDFS**
- **YARN Timeline Service v.2**
- **Shell script rewrite**
- **Shaded client jars**
- **Intra-data node balancer**

A detailed description can be found the official apache Hadoop release page  
<https://hadoop.apache.org/docs/r3.0.0/>

# No need to save 3 replicas of the data

- Data was replicated 3 times in earlier version of Hadoop i.e. Hadoop version 1.x and 2.x
- The idea of replication this was to ensure data backup in the event of data node failure
- The drawback of this is that it reduces the storage space in the cluster by  $1/3^{\text{rd}}$
- Data need not be replicated 3 times, instead a mechanism called “erasure coding” is used to re-construct the lost data block using parity bits

# Hadoop version 3 uses “erasure coding”

- Lost data can be reconstructed using parity bits concept
- This relieves a maximum of 50% extra storage space in the cluster

# Understanding parity bits encoding

- Let's consider the following data in binary format 110010010101
- Divide the data into 3 parts (blocks) and name it as A, B and C

A : 1100

B : 1001

C : 0101

- Assume each block is of size 128MB
- 3 blocks stored as is would occupy  $128 * 3 = 384\text{MB}$
- On replicating it 3 times for failover recovery , it occupies = 1152 MB

# Understanding XOR logic

A	B	Output
1	1	0
0	0	0
1	0	1
0	1	1

# Calculating the parity block

**A : 1100**

**B : 1001**

**C : 0101**

Block A	Block B	Parity Block (using XOR logic)	<b>Block A Is lost</b>	Block B	Parity Block (using XOR logic)
1	1	0	X	1	0
1	0	1	X	0	1
0	0	0	X	0	0
0	1	1	X	1	1

**The lost bits in block A can be reconstructed from the parity block**



# Reconstructing the lost data the parity block

**A : 1100**

Original contents of block A

A	B	Output
1	1	0
0	0	0
1	0	1
0	1	1



**Block A  
Is lost**

Block B      Parity Block  
(using XOR logic)

X	1	0
X	0	1
X	0	0
X	1	1

**Recall XOR ....**

Block B      Parity Block      **Applying XOR**


1	0	1
0	1	1
0	0	0
1	1	0

**Similarly the block B data can also reconstructed if its lost as long as we have the other data block and parity intact**

This file is meant for personal use by maneel.chauhan@coforge.com only.

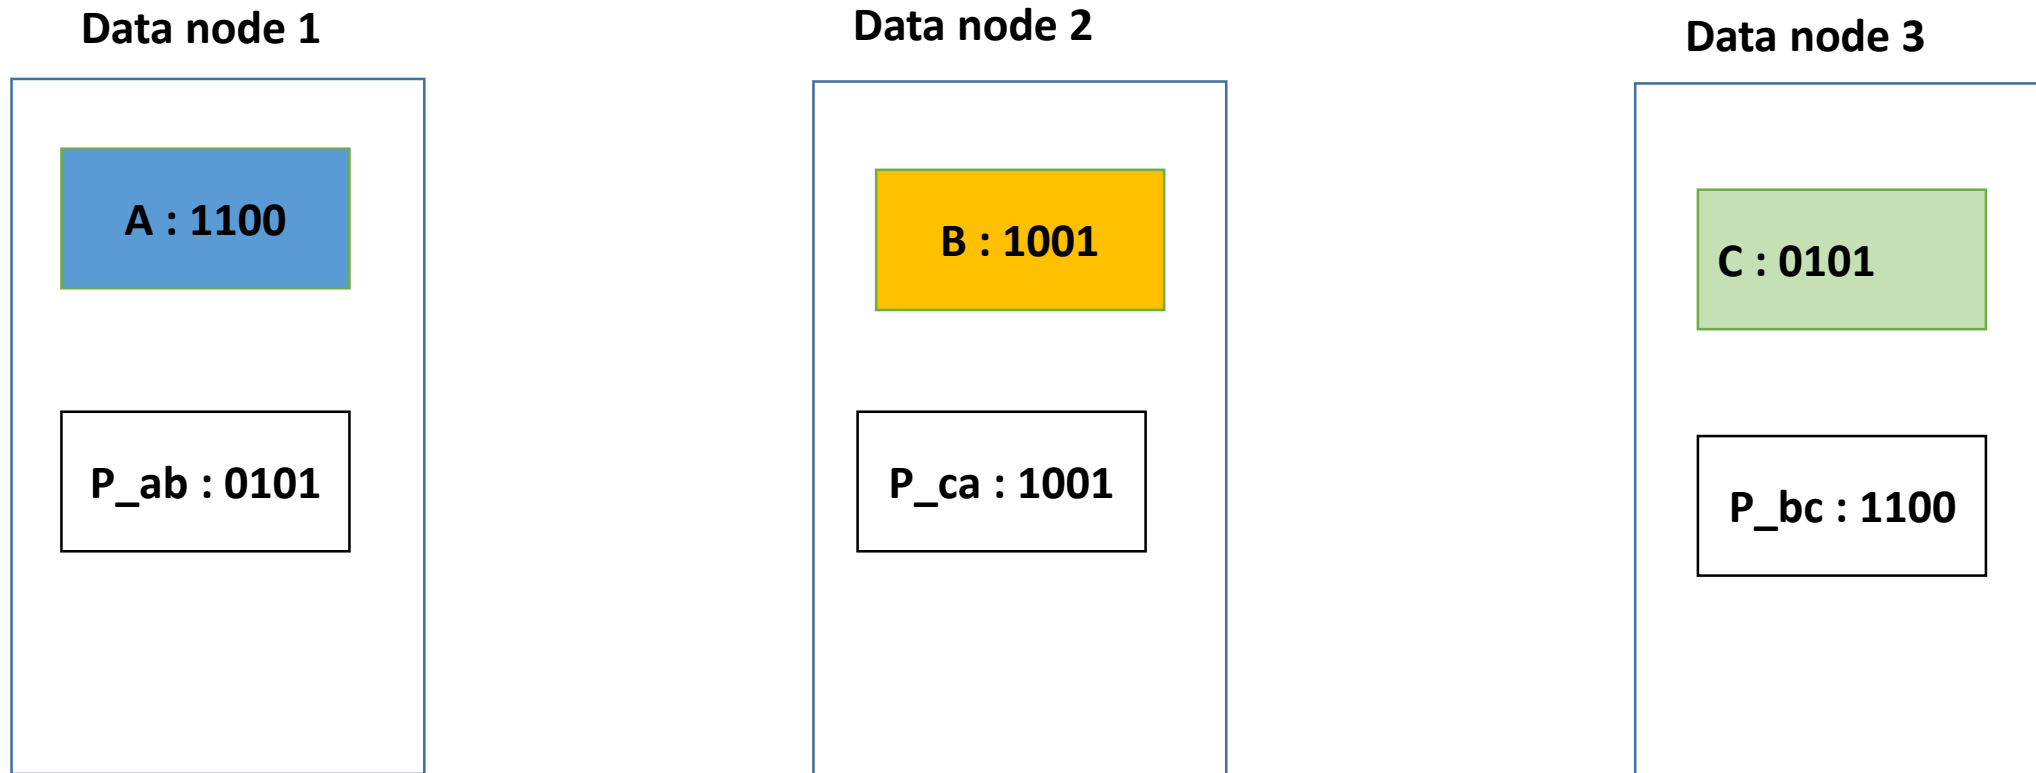
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Saving the data & the parity blocks



- We now have 5 blocks of information eventually (3 data blocks and 2 parity blocks)
- Each block occupies 128MB of data on the disc of a data node
- The 3 data blocks would occupy  $128 * 3 = 384$  and replicating it 3 times occupies 1152 MB
- Saving the above 5 blocks would occupy  $128 * 5 = 640$ MB (almost 40% reduction in disc space)
- Parity blocks are in fact replicated directly or indirectly which would still give up to 30% space saving
- Any savings in disc space would directly reflect on the hardware cost savings

# Saving the data blocks and the parity blocks




**Note:** Erasure coding uses a more complex technique to handle failure recovery using this parity bit concept. The discussion here is limited to basic understanding of block reconstruction.

# Intra data node balancer

- Earlier versions of Hadoop had an inter data node balancer (HDFS balancer)
- This ensured that the data in the cluster was always evenly distributed among the nodes
- In case we had 10GB of data and 5 data nodes, each data node would get 2GB of data
- Typically each data node in the Hadoop cluster has several hard discs
- Intra data node balancer ensures that data is evenly distributed among these nodes

# HDFS balancer

- HDFS balancer ensures that the data is evenly distributed across the data nodes in a cluster
- An example, 12TB of data would be distributed as 3 TB on each of the four data nodes
- HDFS balancer would ensure deleting existing data or adding new data would still result in a balanced cluster




This file is meant for personal use by maneel.chauhan@coforge.com only.

Sharing or publishing the contents in part or full is liable for legal action.

# Drawbacks of HDFS balancer

- A data node has multiple hard drives, usually in the order of 10 to 20 separate hard drives each of capacity anywhere between 1TB to 4TB
- HDFS balancer does not ensure if the data in all the hard drives is evenly distributed
- The skewed data distribution in the hard drives of a data node is due to deletion of data and addition of new data into HDFS

**Data node with multiple hard drives**




This file is meant for personal use by maneel.chauhan@coforce.com only.

Proprietary content. All Rights Reserved. Unauthorized use or distribution prohibited.

# Intra data node balancer

- This component ensures that the data is almost uniformly distributed among the all the discs of a data node
- This results in an improved data read performance

**Data node with multiple hard drives**



This file is meant for personal use by maneel.chauhan@coforge.com only.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Summary

- Understand the 2 important features of Hadoop version 3
- Basics of erasure coding
- Intra data node balancer